

ZATO LABS - PRODUCT MANAGER CASE STUDY

TITAN

Why review work leaves the platform - and where AI actually helps



Prepared by Tejas Pandit

15-minute presentation - 5-minute discussion

September 2026

THE PROBLEM I'M SOLVING

Why does 60% of review work leave Zato's intended workflow - and what should we do about it, including AI?

Reviewer Cockpit

Notifications

Agentic AI mandate

Three signals were handed to me. None of them is yet a diagnosis - that's the work of this presentation.

WHAT THE DATA SAYS - AND WHAT IT DOESN'T

40%

of completed work items are reviewed
inside Reviewer Cockpit

60%

are approved via email or spreadsheets,
outside the platform

+22%

QoQ rise in notification-related support
tickets - split ~evenly between “too many”
and “I missed something important”

Leadership wants agentic AI in Titan this year - but hasn't decided where it creates the most value.


SIGNALS ≠ DIAGNOSIS

40% Cockpit usage isn't automatically bad. “Too many” and “missed important” aren't automatically opposites. I'm treating all three as evidence to interrogate, not problems to immediately fix.

HOW INFORMATION ACTUALLY FLOWS THROUGH ZATO

Data comes in cleanly. The gap opens at the very last step.

I want to show I understand what Zato actually does today - the fix has to fit the real product, not an assumed one.



WHAT I'D INVESTIGATE FIRST

Five specific questions, not generic discovery

1

**Adoption or
capability gap?**

Different fixes either way

2

**Internal or
client bypass?**

Different owners entirely

3

**Notifications
severity-blind?**

Same urgency treatment?

4

**How long do
queries block?**

Thinking time vs. forgotten

5

**Does work return
to Zato?**

Audit risk vs. cosmetic

Question 4 carries the most weight - it's the only one tied to a hard, measurable constraint rather than a guess.

WORKING HYPOTHESIS

Building on those five questions - two distinct hypotheses

On Cockpit bypass

More likely an adoption or narrow-capability gap than a missing feature - the workflow already has hard gates against open queries and review points. Something specific is pushing people around it.



On notifications

Not a volume problem - a severity-blindness problem. Everything inside one category likely fires the same way regardless of true urgency, so people feel overwhelmed and still miss the one that mattered.



Both point to one gap: Zato has the data to know what matters, but isn't yet using it to help people act differently. That's where agentic AI has a real job - next slide.

WHERE AGENTIC AI COULD ACTUALLY HELP

Three candidates, in the order I'd invest

1

Pre-submission readiness nudge

Flags likely reviewer pushback before
submission

Upstream fix

2

Query triage & draft response

Triages open queries, drafts response for
approval

Ties to hard gate

3

Cross-job Work Queue

Portfolio-wide “what needs you today” view

Biggest bet - Q2 only

TWO-QUARTER ROADMAP

Q1: diagnose and prototype the safest bet. Q2: earn the rest.

Q1

Cockpit: adoption vs. capability, segment the 60%

Notifications: test the severity-blindness hypothesis

AI: prototype Priority 1 (pre-submission nudge)

if
validated
→

Q2

conditional

Query still the bottleneck → build Priority 2

Cross-job gap validated → invest in Priority 3

Neither validated → build neither, revisit instead

SUCCESS METRICS

Tied to the actual bottleneck, not generic PM metrics

Outcome

% of eligible reviews completed end-to-end in Zato, without an external approval step

- not Cockpit logins, the actual completed action

Query resolution

Median time an open query blocks job progression

- directly tied to the one hard, provable gate in the workflow

Attention quality

% of high-priority items acted on within expected SLA

- tests whether severity-blindness is actually improving

Guardrail: missed/overdue high-priority work and notification-related tickets must not rise while we ship this.

THE QUICK-FIX DECISION

“2 weeks” vs. “a full quarter” - how I'd actually decide

Does the quick fix meaningfully reduce customer pain, without creating debt that blocks the real fix?



YES - ship it now

Ship as an explicit, named stopgap alongside the quarter-long fix, not instead of it. Tell stakeholders: “this buys relief now; the real fix still ships next quarter as planned.”



NO - don't ship it

If it only hides the symptom or adds complexity the real fix will have to unwind, I wouldn't spend the capacity. I'd be transparent about why waiting is the faster path overall.

WHERE I'D GO FROM HERE

Titan should evolve from a system that tells users what happened, into one that helps them understand what matters and act on it.

What I don't know yet - and would validate before committing further:

- Whether Cockpit bypass is really adoption or a specific, nameable capability gap
- Whether severity-blindness actually explains the missed-important tickets
- Whether query-blocking time justifies Priority 2 before the cross-job hypothesis is tested

Thank you - happy to go deeper on any of this.

